

Methods for handling deadlock

1. Deadlock prevention or Avoidance :-

The idea is not let the s/m into deadlock state.

2. Deadlock Detection & Recovery :-

Let deadlock occur, then do preemption to handle it once occurred.

3. Ignore the pblm all together :-

If deadlock is very rare, then let it happen & reboot the s/m. This is the approach that both Windows and UNIX take.

— * —

1. Deadlock Prevention or Avoidance

* Provides a set of methods to ensure that atleast one of the necessary conditions cannot hold.

• Mutual Exclusion :-

→ To violate this condition, all the resources in the s/m should be in shareable mode.

→ But in s/m, there are some resources are non-shareable in nature.

→ So mutual exclusion must hold.

• Hold and Wait :-

→ One protocol that we can use requires each process to request & be allocated all its resources before it begins execution.

→ Alternative protocol allows a process to request resources only when it has none.

→ These 2 protocols has 2 main disadvantages :-

- Resource utilization may be low, since resource may be allocated but unused for a long period.

- Starvation is possible :- A process that needs several resources needs to wait indefinitely.

• No Pre-emption :-

→ If a process is holding some resources & requests another resource that cannot be immediately allocated to it then all resources, the process is currently holding are preempted.

→ The process will be restarted only when it can regain its old resources, as well as the new ones that is requesting.

→ Alternatively, if a process requests some resources, we first check whether they are available.

→ If they are, we allocate them.

→ If they are not, we check whether they are allocated to some other processes that is waiting for additional resources.

→ If so we preempt the desired resources from waiting processes and allocate them to the requesting process.

→ If the resources are neither available nor held by waiting process, the requesting process must wait.

→ A process can be restarted only when it is allocated new resources it is requesting & recovers all resources that were preempted while it was waiting.

• Circular Wait :-

→ A natural no: is assigned to every resource.

→ Each process is allowed to request for the resources either in only increasing or only decreasing order of the resource number.

→ In case increasing order is followed, if a process requires a lesser number resource, then it must release all the resources having larger number and vice-versa.

→ This approach is the most practical approach & implementable.

→ However, this approach may cause starvation but will never lead to deadlock.

Eg:-

R1 R2 R3 R4 R5

$P_1 \rightarrow R_2, R_4, R_5$ / R_5, R_4, R_2

If P_1 needs R_3 , then R_4 & R_5 are released & remain R_2 . Then R_3 is allocated, then R_4, R_5 .

Deadlock Avoidance / Prevention Methods

- * This strategy involves maintaining a set of data using which a decision is made whether to entertain the new request or not.
- * If entertaining the new request causes the s/m to move in an unsafe state, then it is discarded.
- * This strategy requires that every process declares its maximum requirement of each resource type in the beginning.
- * The main challenge with this approach is predicting the requirement of the processes before execution.
- * Banker's Algorithm is an example of a deadlock avoidance strategy.

Banker's Algorithm

* To implement Banker's Algorithm, follo: 4 data structures are used :-

1. Available - Single Dimensional Array. (i/p)
2. Max - Two Dimensional Array. (i/p)
3. Allocation - Two Dimensional Array (i/p)
= Max - Allocated (Two-Dimensional Array) (o/p)
4. Need

Safe State Sequence \rightarrow o/p.

Process	Allocation			Max			Available			Need		
	A	B	C	A	B	C	A	B	C	A	B	C
P ₁	0	1	0	7	5	3	3	3	2	7	4	3
P ₂	2	0	0	3	2	2				1	2	2
P ₃	3	0	2	9	0	2				6	0	0
P ₄	2	1	1	2	2	2				0	1	1
P ₅	0	0	2	4	3	3				4	3	1

Work = Available

Work		
A	B	C
3	3	2
2	0	0
5	3	2
2	1	1
7	4	3
0	0	2
7	5	5
0	0	0

Finish				
P ₁	P ₂	P ₃	P ₄	P ₅
F	F	F	F	F
	T		T	
				T
T		T		

Check Need \leq Available $\begin{matrix} 7, 5, 5 \\ 3, 0, 2 \\ \hline 10, 5, 7 \end{matrix}$

(1) P₁ \rightarrow 7, 4, 3 \leq 3, 3, 2 \times

(2) P₂ \rightarrow 1, 2, 2 \leq 3, 3, 2 \checkmark

$$\begin{array}{r} 3, 3, 2 \\ + \\ 2, 0, 0 \\ \hline 5, 3, 2 \end{array} \rightarrow \text{Allocation } P_2$$

Update available

(6) P₁ \rightarrow 7, 4, 3 \leq 7, 4, 5 \checkmark

$$\begin{array}{r} 7, 4, 5 \\ + \\ 0, 1, 0 \\ \hline 7, 5, 5 \end{array} \text{ Allocation of } P_1$$

7, 5, 5

(7) P₃ \rightarrow 6, 0, 0 \leq 7, 5, 5 \checkmark

$$\begin{array}{r} 7, 5, 5 \\ + \\ 3, 0, 2 \\ \hline 10, 5, 7 \end{array} \text{ Allocation of } P_3$$

10, 5, 7

(3) P₃ \rightarrow 6, 0, 0 \leq 5, 3, 2 \times

(4) P₄ \rightarrow 0, 1, 1 \leq 5, 3, 2 \checkmark

$$\begin{array}{r} 5, 3, 2 \\ + \\ 2, 1, 1 \\ \hline 7, 4, 3 \end{array} \rightarrow \text{Allocation } P_4$$

(5) P₅ \rightarrow 4, 3, 1 \leq 7, 4, 3 \checkmark

$$\begin{array}{r} 7, 4, 3 \\ + \\ 0, 0, 2 \\ \hline 7, 4, 5 \end{array} \text{ Allocation of } P_5$$

7, 4, 5

\therefore Safe State // $\langle P_2, P_4, P_5, P_1, P_3 \rangle$

Safety Algorithm

- (1) Let $Work$ & $Finish$ be vectors of length ' m ' & ' n ' respectively.
 - Initialize : $Work = Available$.
 - $Finish[i] = false$; For $i=1, 2, 3, 4, \dots, n$
- (2) Find an i s. both
 - $Finish[i] = false$.
- (3) $Need_i \leq Work$ if no such i exists goto step (4).
 - $Work = Work + Allocation[i]$
 - $Finish[i] = true$.
 - goto step (2)
- (4) if $Finish[i] = true$ for all i .
 - then the s/m is in safe state.

Resource Request Algorithm

- 1) If $Request_i \leq Need_i$ P_i Req. \times Need
Eg: $10, 5, 2 \leq 7, 4, 3$
Goto step (2) ; o.w, raise an error condition, since the process has exceeded its maximum claim.
- 2) If $Request_i \leq Available$ \times
Eg: $4, 3, 2 \leq 3, 3, 2$
Goto step (3) ; o.w, P_i must wait, since the resources are not available.
- 3) Have the s/m pretend to have allocated the requested resources to process P_i by modifying the state as follows:-

$$\text{Available} = \text{Available} - \text{Request}_i$$

$$\text{Allocation}_i = \text{Allocation}_i + \text{Request}_i$$

$$\text{Need}_i = \text{Need}_i - \text{Request}_i$$

$$\text{Eg: } P_1 \text{ req: } \begin{matrix} 3 & 2 & 1 \end{matrix} \leq \begin{matrix} 7 & 4 & 3 \end{matrix} \text{ Need.}$$

$$\leq \begin{matrix} 3 & 3 & 2 \end{matrix} \text{ Available}$$

$$\therefore \text{Available} = \begin{matrix} 3 & 3 & 2 \\ 3 & 2 & 1 \\ \hline 0 & 1 & 1 \end{matrix}$$

$$\text{Allocation} = \begin{matrix} 0 & 1 & 0 \\ 3 & 2 & 1 \\ \hline 3 & 3 & 1 \end{matrix}$$

$$\text{Need} = \begin{matrix} 7 & 4 & 3 \\ 3 & 2 & 1 \\ \hline 4 & 2 & 2 \end{matrix}$$

2. Deadlock Detection & Recovery

* When a deadlock detection Algorithm determines that a deadlock has occurred in the s/m, the s/m must recover from that deadlock..

* 2 approaches of breaking a deadlock:-

→ Process Termination

→ Resource Pre-emption.

→ Process Termination :- To eliminate the deadlock, we can simply kill one or more processes. For this, we use 2 methods:

✓ Abort all the deadlocked Processes :

1. Aborting all the processes will certainly break the deadlock, but with a great expenses.
2. The deadlocked processes may have computed for a long time & the result of those partial computations must be discarded & there is a probability to recalculate them later.

Eg: process pid : 5906 Kill -9 5906

✓ Abort one process at a time until deadlock is eliminated :

1. Abort one deadlocked process at a time, until deadlock cycle is eliminated from the s/m.
2. Due to this method, there may be considerable overhead, because after aborting each process, we must run deadlock detection algorithm to check whether any processes are still deadlocked.

→ Resource Pre-emption :- To eliminate deadlocks using resource preemption, we pre-empt some resources from processes & give those resources to other processes. This method will raise 3 issues :-

✓ Selecting a Victim :

We must determine which resources and which processes are to be preempted & also the order to minimize the cost.

✓ Roll Back:

1. We must determine what should be done with the process from which resources are pre-empted. One simple idea is total rollback.
2. That means abort the process & restart it.

✓ Starvation:

1. In a s/m, it may happen that same process is always picked as a victim.
2. As a result, that process will never complete its designated task.
3. This situation is called Starvation & must be avoided.
4. One solution is that a process must be picked as a victim only a finite number of times.