



KTU
NOTES
The learning companion.

**KTU STUDY MATERIALS | SYLLABUS | LIVE
NOTIFICATIONS | SOLVED QUESTION PAPERS**



Website: www.ktunotes.in

Process Synchronization

Producer Process

reg ₁ = counter
reg ₁ = reg ₁ + 1
counter = reg ₁

while (1) {

 /* produce an item in next produced */

 while (counter == (BUFFER_SIZE))

 ; /* do nothing */

 buffer[in] = nextProduced;

 in = (in+1) % BUFFER_SIZE;

 counter++;

Consumer Process

reg ₂ = counter
reg ₂ = reg ₂ - 1
counter = reg ₂

while (1) {

 while (counter == 0)

 ; /* do nothing */

 nextConsumed = buffer[out];

 out = (out+1) % BUFFER_SIZE;

 counter--;

 /* consume the item in next consumed */

at last $out = 5$, $counter = 4$, $reg_2 = 4$

P₀ : producer execute $reg_1 \leftarrow counter$ $\{reg_1 = 5\}$

P₁ : producer execute $reg_1 \leftarrow reg_1 + 1$ $\{reg_1 = 6\}$

P₂ : consumer execute $reg_2 \leftarrow counter$ $\{reg_2 = 5\}$

P₃ : consumer execute $reg_2 \leftarrow reg_2 - 1$ $\{reg_2 = 4\}$

P₄ : producer execute $counter \leftarrow reg_1$ $\{counter = 6\}$

P₅ : consumer execute $counter \leftarrow reg_2$ $\{counter = 4\}$

- It is an incorrect state, as $counter = 4$ as there are 4 full buffer but in fact, there are 5 full buffers.
- An incorrect state arose because, both process manipulate the variable counter concurrently.

A situation where several processes access & manipulate same data concurrently & outcome depends on a particular order in which access takes place is called race condition.

To avoid race condition, it must be ensured that only one process at a time can be allowed to manipulate the variable counter.

Critical Section Problem.

- Consider a system consisting of n processes $\{P_0, P_1, \dots, P_{n-1}\}$
- Each process has a segment of code called critical section, in which the process may be:
 - ✓ changing common variable
 - ✓ updating a table
 - ✓ writing a file
- When one process is executing in its critical section, no other process is to be allowed to execute in its CS.
- Each process must request permission to enter its CS. The section of code implementing this request is the entry section. The CS may be followed by an exit section. The remaining code is the remainder section.

entry section

critical section

exit section

remainder section

A solution to CS problem must satisfy: (5)

✓ Mutual Exclusion:

- If a process P_i is executing in CS, then no other process can be executing in their CS.

✓ Progress Waiting:

- If no other process is executing in its CS & if some other process is waiting then the process that is not in the remainder section can enter into the critical section.

✓ Bounded Waiting:

- There exists a bound on the no. of times that other processes are allowed to enter its CS after a process has made a request to enter its CS & before that request is granted.

Two Process Solution:

- Solution to CS problem is obtained - PETERSON PROBLEM
- Restricted to two processes - Problem

SOLUTION
PROBLEM
↳ Software Solution

✓ boolean flag[2];

✓ int turn;

Initially $\text{flag}[0] = \text{false}$, $\text{flag}[1] = \text{false}$ & turn can either be 0 or 1.

do {

```
    flag[i] = true;
    turn = j;
    while (flag[j] && turn == j);
```

critical section

```
    flag[i] = false;
```

remainder section

}, where i & j are indices of the processes.

To enter the CS, P_i first sets $\text{flag}[i]$ to be true & then sets turn to j , ie, if other process wishes to enter the CS it can do so. If both processes wishes to enter at same time, turn will be set to both i & j at roughly the same time.

The value of turn decides which of the 2 processes is allowed to enter its CS first.

we need to show that:

✓ mutual exclusion is preserved

✓ progress req. is satisfied

✓ bounded waiting req. is met

To prove prop-1: do a state with following val.

- Each P_i enters its CS only if either $\text{flag}[j] == \text{false}$

or $\text{turn} == i$. no serial entries if $\text{turn} \neq i$.

- If both the processes can be executing in their CS at same time, then $\text{flag}[0] == \text{flag}[1] == \text{true}$.

- This indicates that P_0 & P_1 cannot execute while statement at same time, since value of turn can either be 0 or 1.

- Since at a time, $\text{flag}[j] = \text{true}$ & $\text{turn} = j$ & this condition exists as long as P_j is in CS. Thus, mutual exclusion preserved.

To prove prop 2 & 3:

Process P_i can be prevented from entering CS only if it is stuck in while loop with $\text{flag}[j] = \text{true}$ & $\text{turn} = j$. If P_j is not ready to enter CS, then $\text{flag}[j] = \text{false}$ & if P_j has set $\text{flag}[j] = \text{true}$ & is also executing in while statement, then either $\text{turn} = i$ or $\text{turn} = j$.

If $\text{turn} = i$, then P_i will enter CS.

If $\text{turn} = j$, then P_j will enter CS.

Once P_j exits its CS, it will reset $\text{flag}[j] = \text{false}$, allowing P_i to enter its CS.

If P_j resets $\text{flag}[j] = \text{true}$, it must also set $\text{turn} = i$.

Since P_i does not change the value of turn while

executing while statement, P_i will enter the

CS (progress) after atmost one entry by

P_j (bounded waiting).

PETERSON SOLUTION.

turn = 0, 1.

flag = TRUE, FALSE
 (1) (0)

If turn = 0, P₀ turns to enter CS.

If " = 1, P₁ "

P _i	P _j
F	F

turn [i/j] only one variable allowed.

Code	P _i	P _j
flag	T	F

Line 1: flag[i] = F; it needs to enter CS.

turn	J
------	---

Line 2: //allowing other process to participate in CS.
 both cond - T, while - T it loops again bc
 if one cond - F, While - F it enters CS.

while (flag[j] && turn = j);

↳ F; P_i enters CS.

Line 4: release lock; set flag[i] = T.

P _i	P _j
F	F

Why, if P_j needs to enter CS.

P _i	P _j
F	T

turn	i
	F

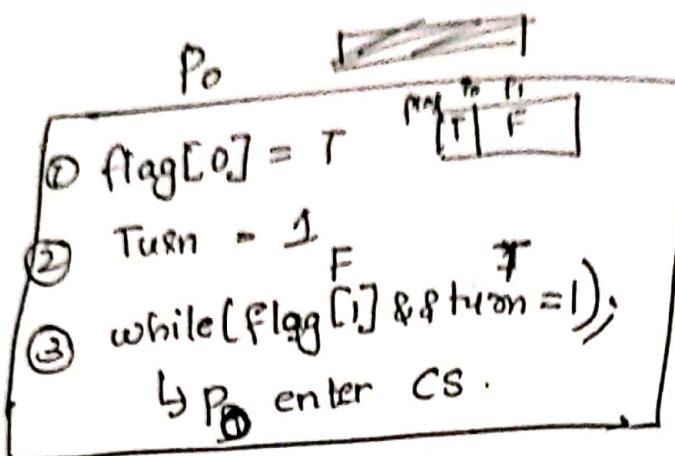
while (flag[i] && turn = i)

↳ F; P_j enters CS

release lock, set flag[i] = T.

Another eg:

	P ₀	P ₁
Flag	F	F
Turn		



	P ₀	P ₁
Flag	T	T

P₁

① flag[1] = T
 ② Turn = 0.
 ③ while(flag[0] & & turn = 0)
 ↳ do not enter CS as P₀ is in CS.

* Req(1) Mutual Exclusion - Satisfied.

* If P₀ release lock only P₁ can enter CS.

i.e., flag[0] = F.

P ₀	P ₁
F	T

→ Thus P₁ can enter CS

* Req(2) - Progress waiting.

P₀ enter CS then RS then sleep.

P₁ enter CS then RS

P₀ doesn't want to enter CS; so it is never going to set flag = T. Thus P₁ can enter CS.

* Req(3) - Bounded waiting.

↳ ^{Process} cannot repeatedly enter into CS if some other process is wanting to enter into CS.

Thus Req(3) failed.

Synchronization Hardware

boolean TestAndSet (boolean &target)
{
 boolean rv = target;
 target = true;
 return rv;
}

(Definition)

If the machine supports TestAndSet instr, then it can implement mutual exclusion by declaring a boolean variable lock, initialized to false.

do {

while (TestAndSet (lock));

critical section .

lock = false;

remainder section .

} while(1);

HARDWARE SYNCHRONIZATION

Limitation of Peterson → Only a process
→ Not adaptable for modern computer arch.

Special H/w instruction

J.

J.
swap()

Test & Set()
Test & set()

- It is an atomic instruction → uninterruptable unit
boolean TestAndSet(boolean & target)

{
 boolean sv = target;

target = TRUE;

return sv;

}

target [T/F]

sv [T/F]

Initially,

target [F]

sv [F]

Line-1
target [F]
sv [F]

Line-2
target [T]
sv [F]

Line-3
sv = F.

initial sv	initial target	sv	target	return sv
F	F	F	T	F
T	F	F	T	F
F	T	T	T	T
T	T	T	T	T

Line-1
target [F]
sv [F]

Line-2
target [T]
sv [F]

Line-3
sv = F

Line-1
target [T]
sv [T]

Line-2
target [T]
sv [T]

Line-3
sv = T

Line-1
target [T]
sv [T]

Line-2
target [T]
sv [T]

Line-3
sv = T

target	return sv
T	T
F	F

do {

- ① while (Test And Set (lock));
- ② CS . set target = lock -

- ③ lock = False;
- ④ RS .

} while (1);

Lock	return av
T	T
F	F

~~lock lock~~

Lock . False

If P_1 needs to enter

Line ① : while (False);

Line ② : ~~&~~ P_1 enter CS

& lock = True

If P_2 need to enter :

Line ① : while (True);

Cannot enter CS.

* Thus Mutual Exclusion Satisfied

when P_1 finishes .

lock = False .

P_2 en checks

Line ① : while (F); & enter CS
& set lock = T .

* Progress .

Initial Lock = F .

, then Lock = T .

P₁
Set Lock = T .
CS .

Lock = F
RS (sleep).

P₂
initial Lock = F
Lock = T .
CS
Lock = F
RS

lock is f , P_1 doesn't want to enter so P_2 can enter as many times as it need .

Thus Progress Waiting Satisfied .

* Bounded waiting :

P₁
 $L = T$
CS
 $L = F$
RS

If P_1 again want to enter CS, it can do as
Lock = F . P_2 ^{should} wait many times
as if P_1 has higher priority .

NO Bounded Limit for P_1 .

- Starvation .

* Bounded waiting Not Satisfied .

If the machine supports swap instrn, then mutual exclusion can be implemented by declaring a global boolean variable lock initialize to false & a local boolean variable key.

```
template<void swap(boolean& a, boolean& b)
```

```
{
```

if (lock == false) {
 lock = true;
 boolean temp = a;
 a = b;
 b = temp;

```
}
```

```
do {
```

```
    key = true;  
    while (key == true)  
        swap(lock, key);
```

Critical section

```
    lock = false;
```

```
} while(1);
```

remainder section

Swap()

```
Void swap(boolean &a, boolean &b)
{
    boolean temp = a;
    a = b;
    b = temp;
}
```

} definition
of swap()

do {

```
key = true;
while (key == true);
swap(lock, key);
```

CS

```
lock = false;
```

R.S.

} while() .

~~Def of swap()~~ Code

Global var. Lock =

Local var. Key =

If key = T will not enter CS

If key = F in Line (3) only it will enter CS

Initially lock = F.

Since Line (1) : key = T.

P₁, P₂ 1 (2) : key == T.

 (3) : swap(lock, key) .

Initially lock = F,

P₁

1: key_{P₁} = T.

2: while(key == T).

3: swap(lock, key)

→ lock = T.

key = F

CS

4: lock = F

P₂

(1) key_{P₂} = T.

(2) Key == T.

(3) swap(lock, key)

↳ True

So it will not
enter CS as

key = T · swap(key = F, lock = T)

key = T, lock = F : Enter CS (as key = F)

Defn of swap

a	b	a	b
T	F	F	T
F	T	T	F
F	F	F	F
T	T	T	T

lock	key	lock	key
F	T	T	F

Thus ME is satisfied

as P₂ is blocked as
key = T. P₁ is in
CS.

If lock = F in P₁ then
only P₂ can enter
CS

<u>Progress</u>	
P ₁	P ₂
CS	CS
lock = F	↓ lock = F
RS	↓ RS.
sleep	

Bounded	Waiting
High Priority P ₁	P ₂
CS	X (1)
lock = F	
RS	X (2)
CS	X (3)
CS	

If P₁ doesn't want to enter CS then it doesn't put lock = F. so P₂ can enter as many times as it need.

P₂ is starving as P₁ is of high priority P₂ is blocked on entering CS.

So bounded waiting cond is not maintained.

(8)

Semaphore

- Semaphore is a synchronization tool.

- Semaphore s is an integer variable that can be accessed only through two operations:

✓ wait

✓ signal

wait(s) adds one to s .

while ($s \leq 0$);

// no-op

$s--;$

}

signal(s) decreases s .

{ if needed $s++$ }

$s++;$ if $s > 0$

}

* Binary Semaphore & Counting Semaphore

$$S = \{0 \text{ or } 1\}$$

$$S = \{n\}$$

Eg: M.E / C.S probm soln

Eg: Resource Mgmt.

wait(S)

{

while ($S \leq 0$); // T → repeat
 // F → dec. S.

$S--;$

}

signal(S)

{

$S++;$

}

1) Semaphore usage

- Semaphore can be used to deal with n-process CS problem.
- The n processes share a semaphore mutex initialized to 1.
- Semaphore can solve synchronization problem.

Example of statement **wait(mutex);** in a process question if not bata about critical section

signal(mutex);

remainder section.

Y while (1);

Consider, two concurrently running process, P₁ with a statement S₁ & P₂ with a statement S₂.

Suppose, S₂ should be executed only after S₁ has completed.

For this, let P₁ & P₂ share a common semaphore

synch, initialized to 0.

The statements, S₁;

signal(synch); in process P₁.

The statement,

wait(synch);

in process P₂.

If s₂; is followed by s₂, in process P₂, will execute s₂ and if synch is initialized to 0, P₂ will execute s₂ only after P₁ has invoked signal(synch), which is only after P₁ has executed s₁.

① SEMAPHORE USAGE

Global var (shared by all)

↳ mutex = 1.

do

{

wait (mutex);

C.S -

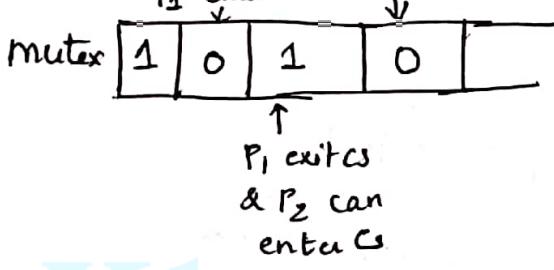
signal (mutex);

R.S

} while (1);

P₁ enters CS

P₂ enters CS



wait (mutex)

{ while (mutex <= 0);

mutex--;

}

signal (mutex)

{ mutex++;

}

P₁
[CS]

P₂

while (0 <= 0); T

(waits till mutex = 1).

P₁
wait (mutex)

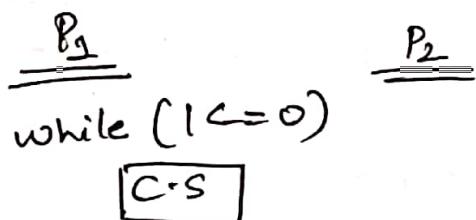
{ while (1 <= 0); F
mutex--;

}

P₁ enters CS; make mutex = 0.
now no other process can enter CS.
P₂ cannot enter as mutex = 0.
when P₁ finishes, it will
signal () & inc. mutex & now
mutex = 1 & P₂ can enter CS.

Thus Mutual Exclusion maintained.

Problems



② P₁ doesn't want to enter CS.
if wont set mutex = 1. So
P₂ can enter CS as many times
as it needs.

Bounded Waiting.

P₁ enters CS → R.S (mutex = 0)

As P₁ (High priority) it repeats the process.

thus P₂ is not allowed to enter CS. i.e., its starving

thus this cond. is not maintained.

SEMAPHORE APPLICATIONS

① CRITICAL SECTION (using mutex)

↳ ORDERING OF STATEMENT

↳ RESOURCE MANAGEMENT

* Ordering of Stmt

P₁

↳ executing stmt S₁

S₁ executing first. S₂ will never execute before S₁. initially

P₂

↳ S₂

$\boxed{S=0 \quad 1 \quad 0}$

S₁; wait(s);

signal(s); S₂;

P₁

S₁;
signal(s)

wait(s); F

s--

(I ≤ 0); F

so s--; (0)

and P₁ enters CS.

~~P₂~~ S₂ shouldn't execute before S₁. So to check

that cond;

~~S=0~~ ir

wait(s);

S=0 initially

while ($S \leq 0$); T so it loops thus ~~#~~ cannot enter CS

So S₂ cannot execute before S₁.

* Resource Management

Resource - Pointer 5

Process - P₁, P₂, P₃, P₄, P₅, P₆ are competing for resource pointer.

Any process that want pointer must perform wait(s)
Once process release pointer it must perform signal(s)

Signal(s)	wait(s)
{	{
s++;	while($sc > 0$)
}	s--;

Initially, $S=5$ (5 resources) - Counting Semaphore

P₁ request pointer,
 wait(s)
 {
 while ($s <= 0$);
 $s--$; }
 } use the pointers

signal(s)
 {
 $s++$;
 } release the pointer

(5)

S	5
P ₁	4
P ₂	3
P ₃	2
P ₄	1
P ₅	0
P ₆	2 X
	1 P ₂ release

Now P₆ comes because it does wait(s); while checking
 while ($s <= 0$); it is T so it will be looping.
 So whenever a resource is not available, that process
 must wait.

If P₂ release pointer, it will perform signal(s)
 so as $s=0$; after signal ($s=1$)

Now P₆ perform wait(s); $\rightarrow (l <= 0)$; F so
 P₆ will use pointer.

2) Synchronization Implementation

- The main disadvantage of semaphores is busy waiting.
- While a process is in its CS, any other process that tries to enter CS must loop continuously in entry code.
- This type of semaphore is called spin lock.
- To overcome the need for busy waiting, modify the definition of wait & signal semaphore operations. When a process executes the wait operation & finds that the semaphore value is not positive, it must wait. Rather than busy waiting, the process can block itself. The block operation places a process into a waiting queue. A process that is blocked, waiting on a semaphore, should be restarted when some other process executes a signal operation. The process is restarted by a wakeup operation, which changes the process from waiting state to ready state. The process is then placed in ready queue.

```
typedef struct
```

```
    {  
        int value;  
        struct process *L;  
    } semaphore;
```

⑦

```
void wait (semaphore s) does the following  
{  
    if (s.value >= 0) {  
        s.value --;  
        if (s.value < 0)  
            {  
                add this process to s.L;  
                block();  
            }  
    }  
    else {  
        signal(s);  
        remove a process P from s.L;  
        wakeup(P);  
    }  
}
```

- $\text{wait}^{(1)}(\text{P}_1)$ trying to enter CS
 $s\cdot \text{value}--; 1 \rightarrow 0$.
 if ($s\cdot \text{value} < 0$) No. it enters
 {
 }
 $\quad // \text{CS (P}_1 \text{ enters CS)}$
- P_2 trying $s\cdot \text{value}--; (0 \rightarrow -1)$
 if ($s\cdot \text{value} < 0$) $-1 < 0$. ✓
 - add P_2 to process list & block.
- P_3 trying $s\cdot \text{value}--; (-1 \rightarrow -2)$
 if ($s\cdot \text{value} < 0$) $-2 < 0$. ✓
 - add P_3 to process list & block.
- P_4 trying $s\cdot \text{value}--; (-2 \rightarrow -3)$
 if ($s\cdot \text{value} < 0$) $-3 < 0$. ✓
 - add P_4 to process list & block
- When P_1 finishes CS it performs $\text{signal}()$
 $s\cdot \text{value}++; -3 \rightarrow -2$.
 if ($s\cdot \text{value} <= 0$)
 remove process from process list (FIFO).
 So P_2 removed from W.Q & it is wakeup.

P_2 enters CS & after performing it call $\text{signal}()$
 $-2 \rightarrow -1$
 $-1 <= 0$ & P_3 wake up.

P_3 enters CS & then $\text{signal}()$
 $-1 \rightarrow 0$.

$0 <= 0$. remove P_3 & P_4 wake up.

We avoided busy-waiting by adding in waiting queue.

S	1	0	-1	-2	-3
---	---	---	----	----	----

-2	-1	0
----	----	---

P ₂
P ₃
P ₄

w.Q

3) Deadlocks & Starvation

- Two or more processes are waiting indefinitely for an event that can be caused only by one of the waiting process. When such a state is reached, the process are in deadlocked.

Two process P₀ & P₁, each accessing two semaphores S & Q, set to value 1:

P ₀ .	(a) Thread	(b) Thread
	: (1) Busy	: (1) Busy
wait(S);	↓	
wait(Q);	↓	wait(Q);
	↓	↓
signal(S);	↓	signal(Q);
signal(Q);	↓	signal(S);

Suppose that P_0 executes $\text{wait}(s)$ & then P_1 executes $\text{wait}(q)$. When P_0 executes $\text{wait}(q)$, it must wait until P_1 executes $\text{signal}(q)$. Similarly, when P_1 executes $\text{wait}(s)$, it must wait until P_0 executes $\text{signal}(s)$. Since these signal operations cannot be executed, P_0 & P_1 are deadlocked.

Another problem related to deadlock is indefinite blocking or starvation, a situation where processes wait indefinitely within the semaphore.

The semaphore is known as counting semaphore, its integer value can range over an unrestricted domain.

The most common form of a semaphore is

Classic Problems of Synchronization

(10)

1.) Bounded Buffer Problem

Assume that it consists of n buffers, each capable of holding one item. The mutex semaphore is initialized to value 1. The empty & full semaphores, count the no. of empty & full buffers.

The semaphore empty is initialized to value n .

The semaphore full is initialized to value 0.

do {

produce an item in next p

wait (empty);

wait (mutex);

add nextp to buffer

signal (mutex);

signal (full);

} while (1);

do {

wait (full);

wait (mutex);

remove an item from buffer to nextc.

signal (mutex);

signal (empty);

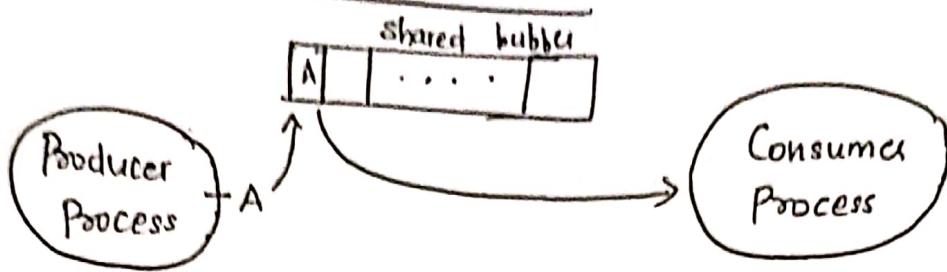
consume the item in nextc

(consume) loop;

} while (1);

CLASSICAL PROBLEMS OF SYNCHRONIZATION

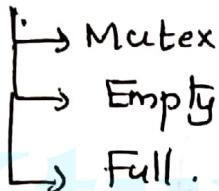
① BOUNDED BUFFER PROBLEM (PRODUCER-CONSUMER PROBLEM)



3 issues :

- * Mutual Exclusion.
- * Overflow - Buffer is full - Producer waits
- * Underflow - Buffer is empty - Consumer waits.

Solution using Semaphore



Structure of producer process

```

do {
    // produce an item
    ...
    wait (empty);
    wait (mutex);
    ...
    // add item to buffer
    ...
    signal (mutex);
    signal (full);
}
while (TRUE);
  
```

```

wait (s)
{
    while (s <= 0);
    s--;
}

// CS.

signal (s)
{
    s++;
}

// RS
  
```

Structure of Consumer Process

do {

// consume an item.

wait(full);

wait(mutex);

.....

// remove an item from buffer

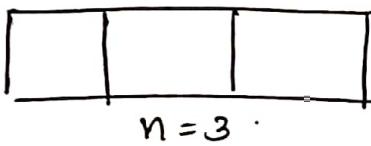
.....

signal(mutex);

signal(full);

} while(TRUE).

① Mutex



mutex
1
0
1
0
1
0
1

Empty
3
2
1
2

Full
0
1
2
1

A	B	L
---	---	---

← consumer consumes A

T	B
---	---

Food-A →

Food B →

~~Food C →~~

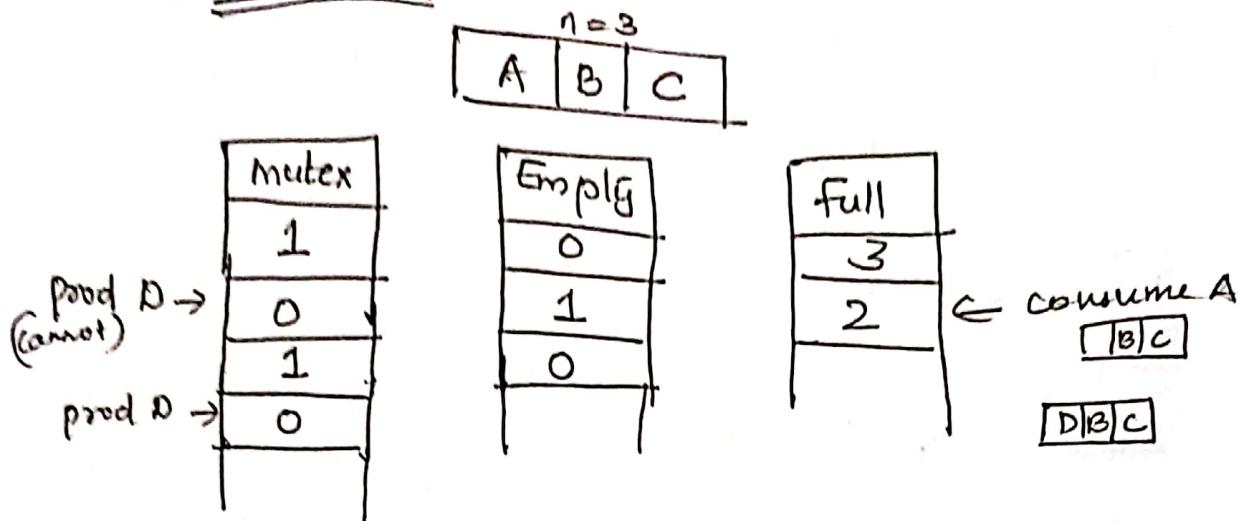
cons. A →

when food is in CS, consumer tries to consume A
wait(full); $1 \rightarrow 0$
wait(mutex); $0 < 0$ (T) doesn't enter CS as producer is in CS. Consumer will be in waiting list.

∴ Mutual Exclusion is maintained.

When producer finishes CS, consumer is moved to CS.

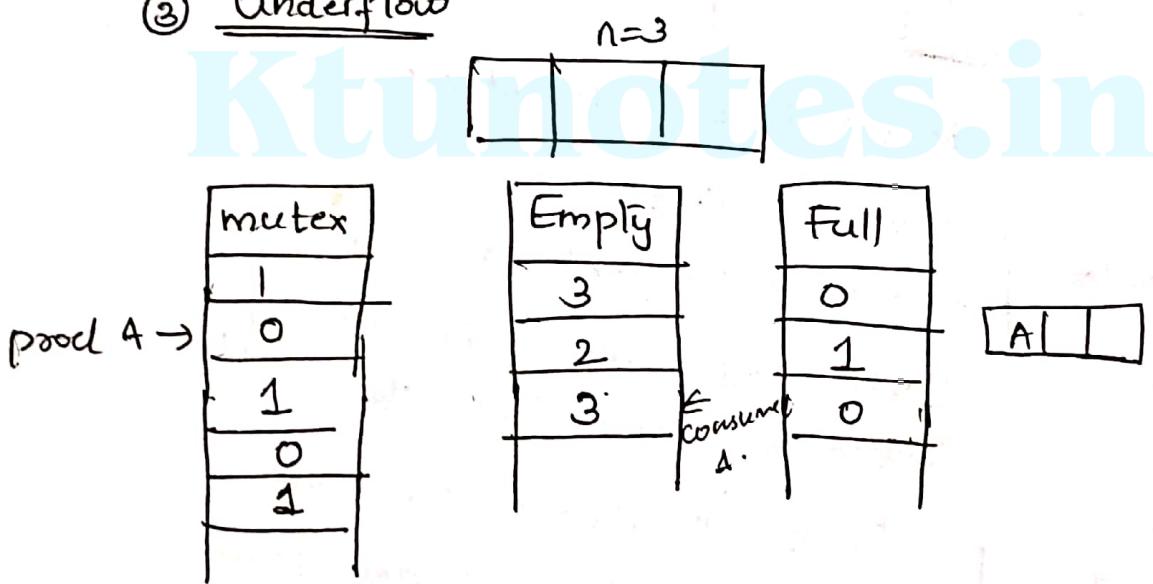
② Overflow



when prod wants to prod D. it is not possible
 $wait(empty); 0 <= 0$ (waiting list).

If thus overflow cond is checked & prod D is pushed into the waiting list when consumer consumes an item only D will be moved to buffer

③ Underflow



when consumer wants to consume, $wait(full), 0 <= 0$
 it cannot consume. It is moved to waiting list
 Thus underflow cond is checked.

If ~~some~~ producer produces an item, then consumer
 can consume that item.

Q.) Readers-Writers Problem

If two readers access the shared data object simultaneously, no adverse effects will result. However, if a writer & some other process access the shared object simultaneously, effects may result. This synchronization problem is referred to as the reader-writer problem.

wait (wrt);
writing is performed
signal (wrt);

Reader process :-

wait (mutex);

readcount++;

if (readcount == 1) signal (wrt);

wait (wrt);

signal (mutex);

reading is performed.

when writer is waiting

readcount--;

if (readcount == 0)

signal (wrt);

signal (mutex);

- The first reader-writer problem, requires that no reader should wait for other readers to finish simply because a writer is waiting.
- The second reader-writer problem requires that if a writing is waiting to access the object, no new readers may start reading.

(11)

A solution to either problem may result in starvation. In the 1st case, writers may starve. In 2nd case, readers may starve.

Data structures,

Semaphore mutex, wait;

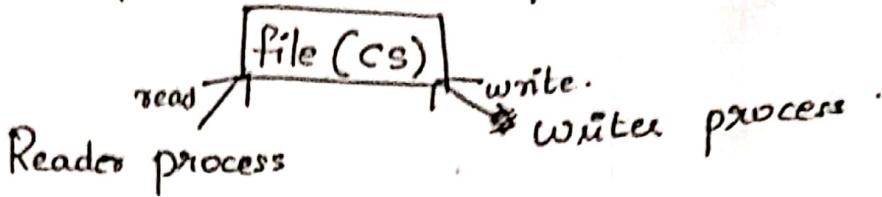
int readcount;

The semaphore mutex & wait is initialized to 1, read count is set to 0. The semaphore wait is common to both reader & writer processes. The mutex semaphore is used when the variable readcount is updated.

②

READERS WRITERS PROBLEM

- * Reader processes & writer processes share a file.



Problems:

1. Allow multiple readers to access the file (R_1, R_2, R_3). (R-R)
2. When reader process is reading the file, no writer process are allowed to access (R-W)
3. When a writer process is writing the file, no reading process or waiting process can access the file. (W-W, W-R)

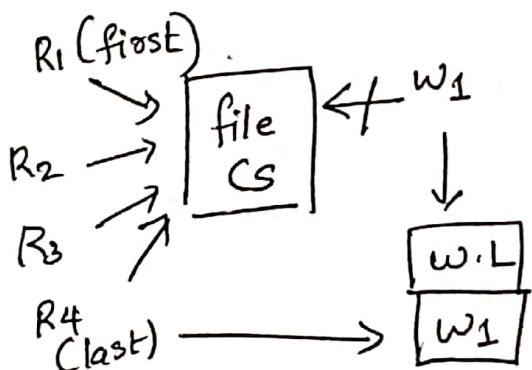
Solution - Semaphores:

1. Mutex = 1 } Semaphore.
2. ~~write~~ = 1.

3. Integer variable readcount = 0.

- * Reader
 - Allow multiple readers.
 - First reader block the writer processes.
 - Last reader allow the waiting writer process to access the file.

- * Writer
 - Blocks all other writer process.
 - Blocks all other reader process.



Structure of a Reader Process

do {

Entry

```
wait(mutex);
readcount++;
if(readcount == 1)
    wait(wrt);
signal(mutex);
```

// Reading is performed
CS

Exit

```
wait(mutex);
readcount--;
if(readcount == 0)
    signal(wrt);
signal(mutex);
```

} while (TRUE);

Structure of a Writer Process

do

Entry

```
wait(wrt);
```

CS //writing is performed

Exit

```
signal(wrt);
```

} while (TRUE);

	mutex	read count	wrt
$R_1 \rightarrow$	1	0	
	0	1	
	1	2	
$R_2 \rightarrow$	0	3	
	1	2	
$R_3 \rightarrow$	0	1	
	1	0	
finished $R_1 \rightarrow$	0		
	1		
finished $R_2 \rightarrow$	0		
	1		
finished last $R_3 \rightarrow$	0		
	1		

w_1 tries $\text{wait}(w_{\text{rt}}) = 0$
 $\leftarrow \text{signal}(w_{\text{rt}})$
 $\leftarrow w_1$
 w_1 then blocks

Thus 3 Conditions of Reader process is satisfied.

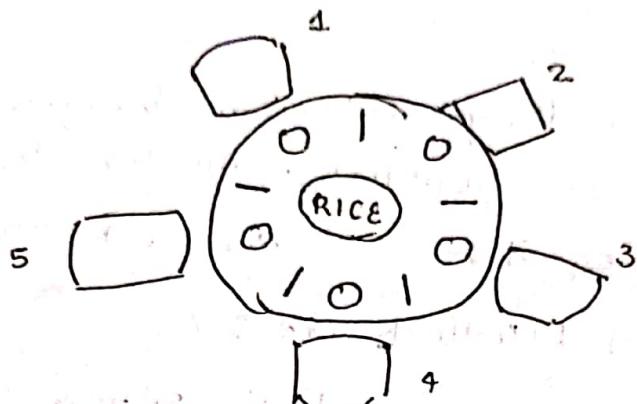
	mutex	read count	wrt
$R_1 \rightarrow$	1	0	
	0	1	

$\leftarrow w_1$

$\leftarrow w_2$ (no access
 $\leftarrow w_1$ can be given as
finished can be given as
 $\text{wait}(w_{\text{rt}})$ fails)

R_1 tries, when performing $\text{wait}(w_{\text{rt}})$ it is blocked.
Thus the 2 conditions of Writer process is satisfied.

Dining-Philosophers Problem .



Consider 5 philosophers who spend their lives thinking & eating. The philosophers share a common circular table surrounded by 5 chairs. In the center of the table is a bowl of rice & table is laid with 5 single chopsticks. A philosopher picks only one chopstick at a time. She cannot pick up a chopstick that is already in the hand of a neighbour. When a hungry philosopher has both her chopsticks at same time, she eats w/o releasing her chopsticks. when she finishes eating, she put down both of her chopsticks & start thinking again. One simple solution is to represent each chopstick by a semaphore.

A philosopher tries to grab the chopstick by executing a wait operation on that semaphore.

She releases her chopstick by executing a signal operation on the semaphore. Shared data are:

semaphore chopstick[5];
where all elements of chopstick are initialized to 1.

do {

wait(chopstick[i]);

wait(chopsticks[(i+1)%5]);

eat

signal(chopstick[i]);

signal(chopstick[(i+1)%5]);

think

} while(i);

Suppose that all 5 philosophers become hungry simultaneously & each grabs her left chopstick. All the elements of chopstick now be equal to 0, a deadlock arises.

Solutions to Avoid Deadlocks are:

1) Allow atmost 4 philosophers to be sitting simultaneously at the table.

2) Allow a philosopher to pick up her chopstick only if both chopsticks are available.

3) An odd philosopher picks up first her left chopstick & then her right chopstick, whereas an even philosopher picks up her right chopstick & then her left chopstick.

DINING PHILOSOPHERS PROBLEM

- * Five philosophers are sitting around a circular table.
- * Dining table has 5 chopsticks & bowl of rice in the middle.
- * Philosophers either eat or think.
- * When philosopher wants to eat, he uses 2 chopsticks.
- * When he wants to think, he keeps down both chopsticks.

* Problems:

Develop an algorithm where no philosopher starves i.e., Every philosopher should eventually get a chance to eat.

Chopstick [5] initialized to 1.

Structure of Philosophers i: P₀

do {

wait (chopstick[i]);

F₀

wait (chopstick[(i+1)%5]);

(i+1)%5 = 1
F₁

// Eat

release locks;

{ signal (chopstick(i));

F₀

signal (chopstick((i+1)%5));

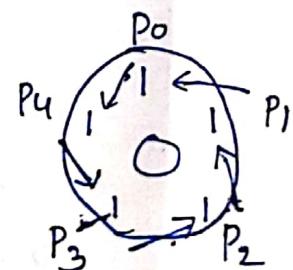
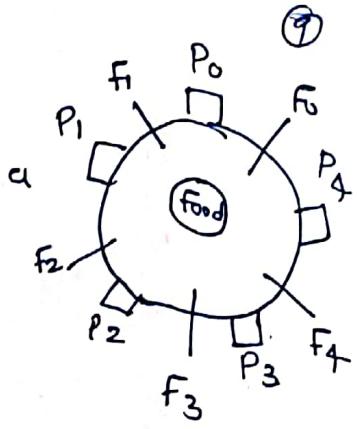
F₁

// Think;

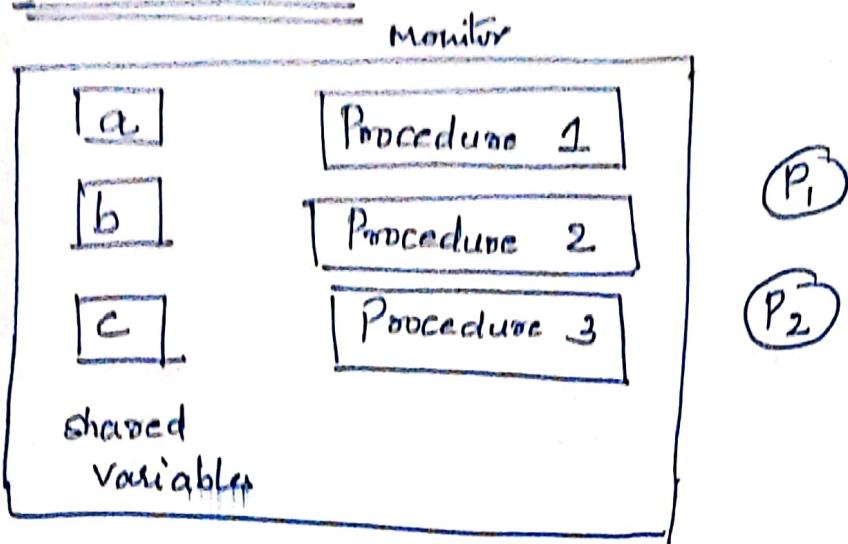
} while (true);

P₀ waits until P₁ releases, P₁ waits for P₂ to release, P₂ waits for P₃ to release.

Everyone is in deadlock state.



MONITORS



P₁ & P₂ has shared variables a,b,c . we need to solve CS problem .

Only one process can enter monitor .

When P₁ enters monitor , P₂ must wait .

P₁ can access the shared variables only through procedures .

If proc-1 is get(a) & proc-2 is set(a) .

Then P₁ must first use proc-1 to get a & P₁ must then use proc-2 to modify value of a .

when P₁ leaves only P₂ can enter monitor .

* Monitor is a module that contains :-

- Shared data .
- Procedures that operate on the shared data

Syntax of Monitors

```

monitor
{
    condition variables; // vars operates using 2 opns:
                           //      wait(), signal()
                           ↓
                           blocked
    variables;
}

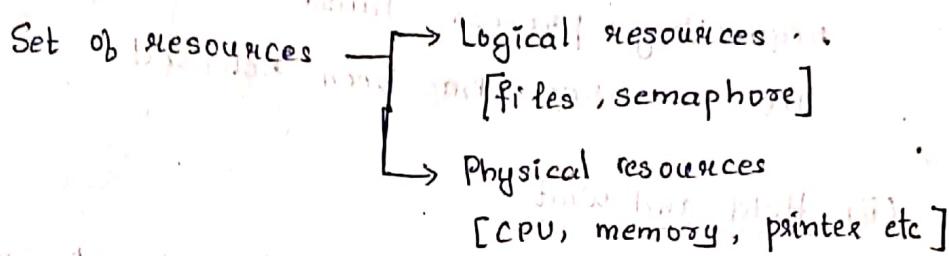
Procedures 1
{
    =
}
Procedures 2
{
    =
}
}

```

DEADLOCK

Process require resource, if not available then moves to waiting state. The resources that are requested are held by waiting process.

System Model



A process must request a resource before using it and must release the resource after using it.

The no. of resources requested may not exceed the total no. of resources available in the system.

Normal mode of operation

1) Request : If the request cannot be granted immediately ie, the resource is being used by another process, then the requesting process must wait until it can acquire the resource.

2) Use : The process can operate on the resource ie, if the resource is a printer, the process can print on the printer.

3) Release : The process releases the resources.

Deadlock Characteristics

4) Necessary Conditions

Deadlock condition hold if all the four conditions hold simultaneously :

(i) Mutual Exclusion

- Atleast one resource must be held in a non-shareable mode ie, if only one process at a time can use the resource.
- If another process requests that resource, the requesting process must be delayed until the resource has been released.

(ii) Hold and wait

- A process must be holding atleast one resource & waiting to acquire additional resources that are currently being held by other processes.

(iii) No Preemption

- Resources cannot be preempted ie, a resource can be released only voluntarily by the process holding it, after that process has completed its task.

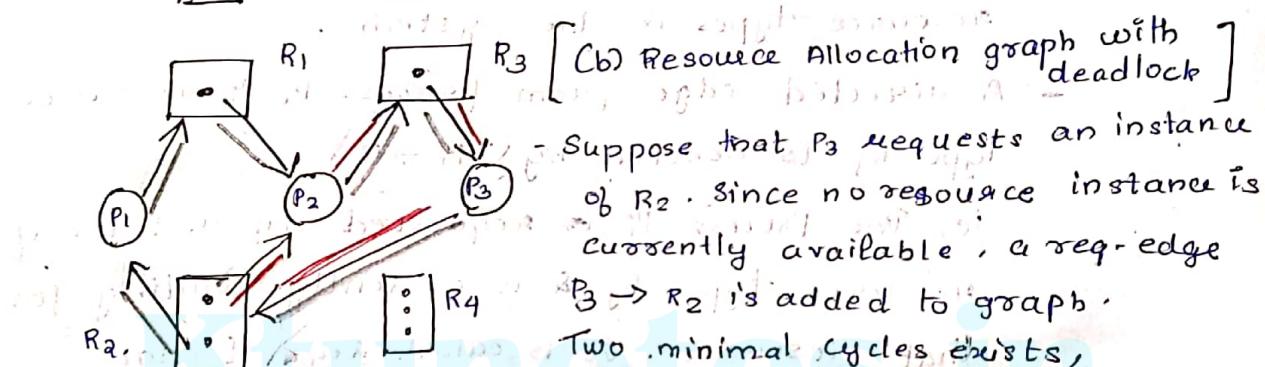
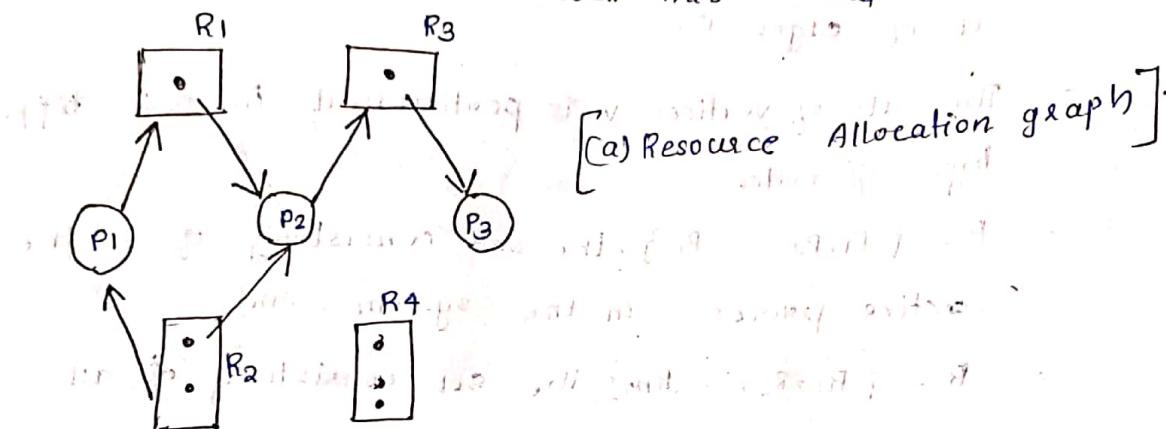
(iv) Circular Wait

- A set $\{P_0, P_1, P_2, \dots, P_n\}$ of waiting process must exist such that P_0 is waiting for a resource that is held by P_1 , P_1 is waiting for a resource that is held by P_2, \dots, P_{n-1} is waiting for a resource that is held by P_n and P_n is waiting for a resource that is held by P_0 .

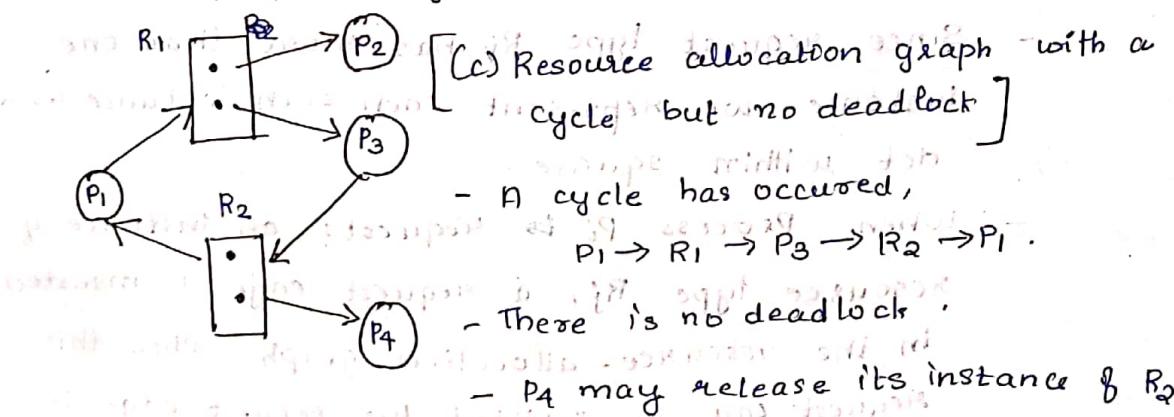
2) Resource Allocation Graph

- Deadlocks can be described in terms of directed graph called a system resource-allocation graph.
- The graph consists of a set of vertices V & a set of edges E .
- The set of vertices V is partitioned into different types of nodes:
 - $P = \{P_1, P_2, \dots, P_n\}$, the set consisting of all the active process in the system.
 - $R = \{R_1, R_2, \dots, R_m\}$, the set consisting of all resource types in the system.
- A directed edge from process P_i to resource type R_j is denoted by $P_i \rightarrow R_j$.
 ie, The Process P_i is requested an instance of resource type R_j & is currently waiting for that resource. It is called request edge.
- A directed edge from resource type R_j to process P_i is denoted by $R_j \rightarrow P_i$.
 ie, An instance of resource type R_j has been allocated to process P_i . It is called Assignment edge.
- Process as circle, resource as square.
- Since request type R_j has more than one instance, we represent each such instance as a dot within square.
- When Process P_i requests an instance of resource type R_j , a request edge is inserted in the resource-allocation graph. When this request can be fulfilled the request edge is transformed to assignment edge. When the process no longer needs access to the resource it releases the resources, & as a result assignment edge is deleted.

- If the graph contains no cycle, then no process in the system is deadlocked.
- If the graph does contain a cycle, then a deadlock may exist.
- If each resource type has exactly one instance, then a cycle implies a deadlock has occurred.



- $P_1 \rightarrow R_1 \rightarrow P_2 \rightarrow R_3 \rightarrow P_3 \rightarrow R_2 \rightarrow P_1$
- $P_2 \rightarrow R_1 \rightarrow P_3 \rightarrow R_2 \rightarrow P_2$
- ✓ P_1, P_2, P_3 are deadlocked.
- ✓ P_2 is waiting for R_3 which is held by P_3 .
- ✓ P_3 is waiting for either P_1 or P_2 to release R_2 .
- ✓ P_1 is waiting for P_2 to release R_1 .



That resource can then be allocated to P_3 , breaking the cycle.

Methods of handling Deadlocks

We can deal with the deadlock problem in one of three ways:

1. Can use a protocol to prevent or avoid deadlock, ensuring that the system will never enter a deadlock state.
2. Can allow the system to enter a deadlock state, detect it & recovery.
3. Can ignore the problem altogether, & pretend the deadlocks never occur in the system.

To ensure that deadlocks never occur, the system can use either a deadlock prevention, nor a deadlock avoidance scheme.

- Deadlock Prevention is a set of methods for ensuring that atleast one of the necessary conditions cannot hold.
- Deadlock Avoidance requires that the OS to give additional information about which resources a process will request & use.

Deadlock Prevention

Ensure that atleast one condition among 4 doesn't hold for the deadlock to occur.

hold:

1. Mutual Exclusion

- Mutual exclusion condition must hold for non-shareable resources.

- Eg: A printer cannot be simultaneously shared by several processes.

- Shareable resources do not require mutually exclusive access & thus cannot be involved in a deadlock.

- If several processes attempt to open a read-only file at the same time, they can be granted simultaneous access to the file.

- A process never needs to wait for a sharable resource.

2. Hold and Wait.

- Whenever a process requests a resource, it does not hold any other resources.

- One protocol that can be used requires each process to request & be allocated all its resources before it begins execution.

- An alternative protocol allows a process to request resource only when the process has none. A process may request some resources & use them. Before it can request any additional resources, however it must release all the resources that it is currently allocated.

We consider a process that copies data from tape drive to a disk files, sort the disk files & then prints the results to a printer. If all the resources must be requested at the beginning of the process, then the process must initially request the tape drive, disk file & printer. It will hold the printer for its entire execution, even though it needs the printer only at the end.

The second method allows the process to request initially only the tape drive & disk file. It copies from the tape drive to disk, then releases both the tape drive & disk file. The process must then again request the disk file & printer. After copying the disk file to the printer, it releases these 2 resources & terminates.

3. No Preemption

$$P_i \rightarrow R_j \rightarrow P_j$$

(7)
(4)

- If a process is holding some resources & requests another resources that cannot be immediately allocated it, i.e., the process must wait, then all resources currently being held are preempted.
- The preempted resources are added to the list of resources for which the process is waiting. The process will be restarted only when it can regain its old resources, as well as the new ones that is requesting.
- Alternatively if a process requests some resources, we first check whether they are available. If they are, we allocate them. $P_i \rightarrow R_j \xrightarrow{?} P_j \rightarrow R_k$. $P_i \rightarrow R_j \checkmark$ then $R_j \rightarrow P_j$.
- If they are not available, we check whether they are allocated to some other process that is waiting for additional resource. If so, we preempt the desired resources from the waiting process & allocate them to the requesting process.
- If the resources are not either available or held by a waiting process, the requesting process must wait while it is waiting, some of its resources may be preempted but only if another process requests them. A process can be restarted only when it is allocated the new resources it is requesting & recovers any resource that were preempted while it was waiting.
- One way to ensure that this condition never holds is to impose a total ordering of all resource types & to require that each process requests resources in an increasing order of enumeration.

4 Circular Wait

One way to ensure that this condition never holds is to impose a total ordering of all resource types & to require that each process requests resources in an increasing order of enumeration.

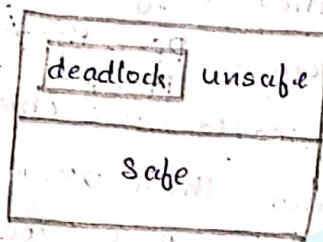
A Function $F: R \rightarrow N$, where N is the set of natural numbers.

$$F(R_0) < F(R_1) < \dots < F(R_n) < F(R_0)$$

Deadlock Avoidance:

- For avoiding deadlock it is to require additional information about how resources have to be requested.
- The simplest and most useful model requires that each process declare the maximum no. of resources of each type that it may need.

4) Safe State:



- A state is safe, if the system can allocate resources to each process in some order & still avoid a deadlock.
- A system is in a safe state only if there is a safe sequence.
- A sequence of processes $\langle P_1, P_2, \dots, P_n \rangle$ is a safe sequence, if for each P_i , the resources that P_i ~~still~~ requests can be satisfied by the currently available resources plus the resources held by all P_j .
- If the resource that process P_i needs are not immediately available, then P_i can wait until all P_j has finished. When they have finished, P_i can obtain all of its needed resources, complete its task, return its allocated resources & terminate. When P_i terminates, P_{i+1} can obtain its needed resource & so on. If no such sequence exists, then system is in unsafe.

- A safe state is not a deadlock state.
- A deadlock state is an unsafe state.
- Not all unsafe states are deadlocks.

Problem:

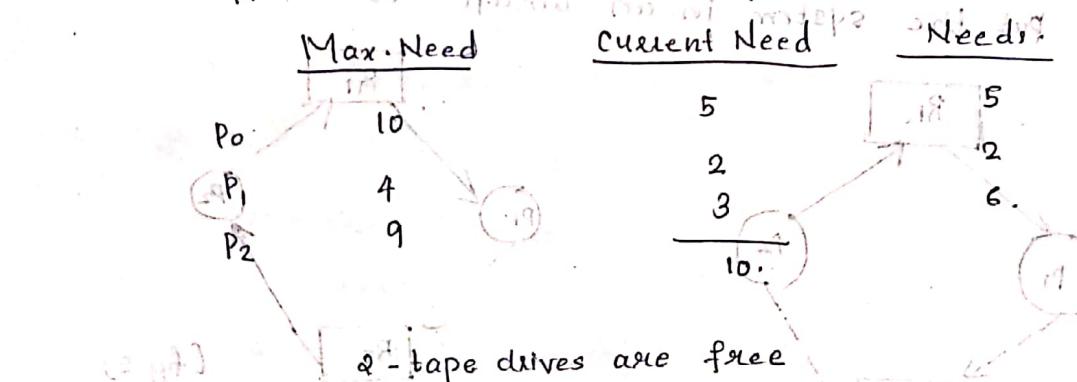
Total : 12 magnetic tape drives.

	<u>Max. Need</u>	<u>Current Need</u>	<u>Needs</u>
P ₀	10	5	5
P ₁	4	2	2
P ₂	9	2	7

all 6 units 3-tape drives are free. & now consider the set $\{P_1, P_0, P_2\}$ - safe state. If P_1 is allocated then when it releases it have 5 available.

Then P_0 can get all its tape drives & return them, now it may have 10 available & finally P_2 could get all 10 & it may return as 12.

Suppose if P_2 requests & is allocated 4 tape drives.



2-tape drives are free.

Here, P_1 can only be allocated all its tape drives.

when it returns, it has 4. Since P_0 needs 5, it cannot satisfy the need of P_0 or P_2 . Thus a deadlock occurs.

After allocation of 4 tape drives, the available tape drives are 8.

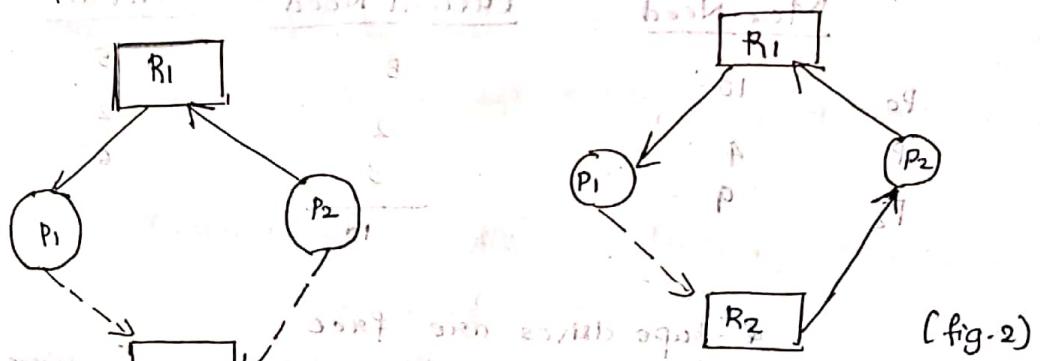
Now, P_1 releases 4 tape drives, so available tape drives are 12.

Process P_2 releases 2 tape drives.

2) Resource Allocation Graph.

A claim edge $P_i \rightarrow R_j$ indicates that process P_i requests resource R_j at some time in future. This edge resembles a request edge in direction, but is represented by dashed lines. When process P_i requests resource R_j , the claim edge $P_i \rightarrow R_j$ is converted to a request edge. Similarly, when a resource R_j is released by P_j , the assignment edge $R_j \rightarrow P_i$ is converted to a claim edge $P_i \rightarrow R_j$. Suppose that process P_i requests resource R_j . The resource can be granted only if converting the request edge $P_i \rightarrow R_j$ into an assignment edge $R_j \rightarrow P_i$ does not result in the formation of a cycle in resource allocation graph.

If no cycle exists, then the allocation of resource will leave the system in a safe state. If a cycle is found, then the allocation will put the system in an unsafe state.



(fig-1). (Resource allocation graph for deadlock avoidance)

[fig-1] - Suppose that P_2 requests R_2 . Although R_2 is currently free, we cannot allocate it to P_2 , since this action will create a cycle in the graph.

[fig-2] - A cycle indicates that the system is unsafe.

If P_1 requests R_2 & P_2 requests R_1 , then a deadlock will occur.

Banker's Algorithm.

↳ Available [j] = k

- k instance of resource type j available.

↳ Max [i,j] = ~~k~~.

- P_i require atmost k instance of resource type R_j.

↳ Allocation [i,j] = k.

- Process P_i is currently allocated k instance of resource type R_j.

↳ Need [i,j]

- Process P_i need k more instance of resource type R_j to complete its task.

$$\text{Need}[i,j] = \text{Max}[i,j] - \text{Allocation}[i,j].$$

Safety Algorithm.

1) Let work & finish be vectors of length m & n resp.

Initialize work = Available &

Finish [i] := false, for i=1, 2, ..., n.

2) Find an i such that both -

(a) Finish [i] = false

(b) Need_i \leq work.

If no such i exists then go to step 4.

3) work = work + Allocation_i.

Finish [i] := true

Go to step 2.

4) If Finish [i] = true for all i, then the system is in a safe state.

Resource Request Algorithm

1) If $\underline{\text{Request}} \leq \underline{\text{Need}_i}$, go to step 2.

Otherwise raise an error condition, since the process has executed its maximum claim.

2) If $\underline{\text{Request}_i} \leq \underline{\text{Available}}$, go to step 3.

Otherwise P_i must wait, since the resources are not available.

3) Have the system pretend to have allocated the requested resource to process P_i by modifying the state as follows:

$$\checkmark \underline{\text{Available}} := \underline{\text{Available}} - \underline{\text{Request}}$$

$$\checkmark \underline{\text{Allocation}} := \underline{\text{Allocation}} + \underline{\text{Request}}$$

$$\checkmark \underline{\text{Need}}_i := \underline{\text{Need}}_i - \underline{\text{Request}}$$

Problem

4)

	<u>Allocation</u>			<u>Max</u>			<u>Available</u>		
	A	B	C	A	B	C	A	B	C
P ₀	0	1	0	7	5	3	3	3	2
P ₁	2	0	0	3	2	2	1	1	1
P ₂	3	0	2	9	0	2	2	2	2
P ₃	2	1	1	2	2	2	3	3	1
P ₄	0	0	2	4	3	3	3	3	2

Soln:

Safety Algorithm

Need [Max - Allocation] Need \leq Avail

A B C ① P₀ 743 \leq 332 X

P₀ 7 4 3 ② P₁ 122 \leq 332 ✓

P₁ 1 2 2 Avail = 332 + 200 = 532

P₂ 6 0 0 ③ P₂ 600 \leq 532 X

P₃ 0 1 1 ④ P₃ 011 \leq 532 ✓

P₄ 4 3 1 Avail = 532 + 211 = 743

⑤ P₄ 431 \leq 743 ✓

Avail = 743 + 002 = 745.

⑥ 902 + 743 \leq 745 ✓

Avail = 745 + 010 = 755.

⑦ P₂ 600 \leq 755 ✓

∴ Avail = 755 + 302 = 1057

∴ A B C \leq 1057

10 5 7

$\therefore \langle P_1, P_3, P_4, P_0, P_2 \rangle$ is safe.

Ques. Request of P₁ comes to 102. Is it safe? ①

Ans. Yes. Request Algorithm. ②

① 102 \leq 122 [Req \leq need]

② 102 \leq 332 [Req \leq Avail]

③ Avail := 332 - 102 = 230 [Avail := Avail - Req]

④ Alloc := 200 + 102 = 302 [Alloc := Alloc + Req]

⑤ Need = 122 - 102 = 020 [Need := Need - Req]

Alloc Max Avail Need

P₁ 302 322 230 020

Part

	<u>Allocation</u>			<u>Max</u>			<u>Available</u>			<u>Need</u>		
	A	B	C	A	B	C	A	B	C	A	B	C
P ₀	0	1	0	7	5	3	2	3	0	7	4	3
P ₁	3	0	2	8	2	2				0	8	0
P ₂	3	0	2	9	0	2				6	0	0
P ₃	2	1	1	2	2	2				0	1	1
P ₄	0	0	2	4	3	3				4	3	1
				8	8	7						

① P₀. 743 \leq 230 X.

② P₁. 020 \leq 230 ✓

$$230 + 302 = 532.$$

③ P₂. 600 \leq 532 X.

④ P₃. 011 \leq 532 ✓

$$532 + 211 = 743.$$

⑤ P₄. 431 \leq 743 ✓

$$743 + 002 = 745.$$

⑥ P₀. 743 \leq 745 ✓

$$745 + 010 = 755.$$

⑦ P₂. 600 \leq 755 ✓

$$755 + 302 = 1057 \quad \therefore \langle P_1, P_3, P_4, P_0, P_2 \rangle \text{ is safe}$$

Ab Req. P₄ (3 3 0) & Req. P₀ (0 2 0) comes.

Req. Algorithm.

① 330 \leq 431 ✓ 330 \leq 230 X. Avail = 230 - 020 = 210

② 020 \leq 743 ✓ 020 \leq 230 ✓ All = 010 + 020 = 030. Need = 743 - 020 = 723.

	<u>Allocation</u>			<u>Need</u>			<u>Avail</u>			
	A	B	C	A	B	C	A	B	C	
P ₀	0	3	0	7	2	3	2	1	0	② P ₁ 020 \leq 210 X
P ₁	3	0	2	0	2	0				③ P ₂ 600 \leq 210 X
P ₂	3	0	2	6	0	0				④ P ₃ 011 \leq 210 X
P ₃	2	1	1	0	1	1				⑤ P ₄ 431 \leq 210 X
P ₄	0	0	2	4	3	1				Undo the changes, since it is unsafe

So change as in Part 1

Deadlock Detection

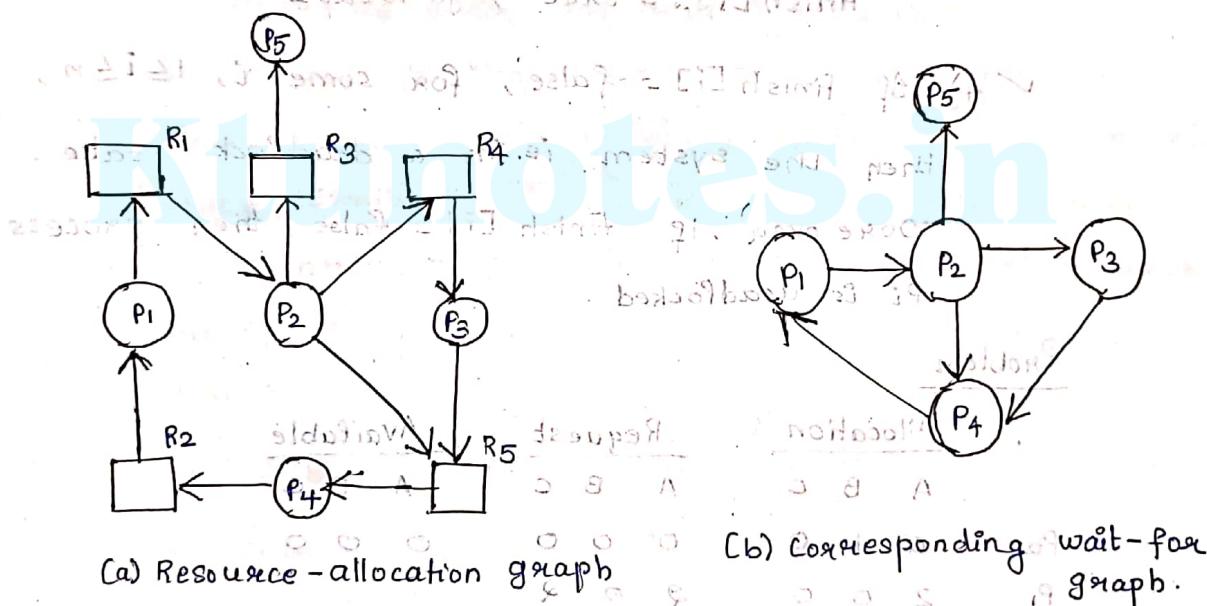
(8)

Determine whether deadlock has occurred.

1) Single Instance of Each Resource Type.

If all the resources have only a single instance, then we can define a deadlock detection algorithm called a wait-for graph.

An Edge from P_i to P_j in a wait-for graph implies that process P_i is waiting for process P_j to release a resource that P_i needs. An edge $P_i \rightarrow P_j$ exists in a wait-for graph if & only if the corresponding resource-allocation graph contains two edges $P_i \rightarrow R_q$ & $R_q \rightarrow P_j$ for some resource R_q .



2.) Several Instances of a Resource Type.

The wait-for graph scheme is not applicable to a resource allocation system with multiple instances of each resource type.

For example:

$$\begin{aligned} & R_1 = 2 \text{ units} \\ & R_2 = 3 \text{ units} \\ & R_3 = 4 \text{ units} \\ & R_4 = 5 \text{ units} \\ & R_5 = 6 \text{ units} \end{aligned}$$

Deadlock detection algorithm

1) Let work & finish be vectors of length m respectively.

Initialize work_i = available, for $i=0, 1, 2, \dots, n$,

If Allocation $\neq 0$, then Finish [i] = false

otherwise

Finish [i] = true

2) Find an index i such that both:

(a) Finish [i] = false

✓ (b) Request_i \leq work_i

If no such i exists, go to step 4

3) work_i = work + Allocation

Finish [i] := true, go to step 2.

✓ 4) If Finish [i] = false, for some i , $1 \leq i \leq n$,

then the system is in a deadlock state.

Moreover, if Finish [i] = false, then process P_i is deadlocked.

Problem -

	Allocation			Request			Available		
	A	B	C	A	B	C	A	B	C
P ₀	0	1	0	0	0	0	0	0	0
P ₁	2	0	0	2	0	2	1	1	1
P ₂	3	0	3	0	0	0	1	1	1
P ₃	2	1	1	1	0	0	1	1	1
P ₄	0	0	2	0	0	2	1	1	1
P ₀	011 \leq avail								
P ₀	000 \leq 000 ✓								
	000 + 010 = 010								
P ₁	202 \leq 010 X								
P ₂	000 \leq 010 ✓								
	010 + 303 = 313								
P ₃	100 \leq 213 ✓								
	213 + 211 = 524								
P ₄	002 \leq 524 ✓								
	524 + 002 = 526								
P ₁	202 \leq 526 ✓								
	526 + 200 = 726								

$\{P_0, P_2, P_3, P_4, P_1\} \rightarrow \text{fin}[i] = \text{true}$

Change $P_2 < 001 >$

P₀ 000 \leq 000 ✓

$$000 + 010 = 010$$

P₁ 202 \leq 010 X

P₂ 001 \leq 010 X

P₃ 100 \leq 010 X

P₄ 002 \leq 010 X

Recovery from Deadlock

1) Process Termination

- ✓ Abort all deadlocked processes.

- This method clearly will break the deadlock cycle, but at a great expense.

- ✓ Abort one process at a time until deadlock cycle is eliminated.

- This method incurs considerable overhead, since after each process is aborted in a deadlock detection algorithm must be invoked to determine whether any processes are still deadlocked.

Aborting a process may not be easy. If the process was in the midst of updating a file, terminating it will leave that file in an incorrect state. Similarly, if the process was in the midst of printing data on the printer, the system must reset the printer to a correct state before printing the next job.

Criteria for process selection :忙iest

- ✓ what the priority of the process is.
- ✓ How long process was computed.
- ✓ How many & what type of resource used.
- ✓ How many more resources needed.
- ✓ How many process will need to be terminated.

Q) Resource Preemption

Issues regarding resource preemption:

✓ Selecting a victim

- Which resources & which processes are to be preempted? As in process termination, we must determine the order of preemption to minimize cost of interrupting a process & breaking it.

✓ Roll back

- If we preempt a resource from a process, then it cannot continue with its normal execution; it is missing some needed resource. We must roll back the process to some safer state & restart it from that state.

✓ Starvation

In a system where victim selection is based primarily on cost factors, it may happen that the same process is always picked as a victim. As a result, this process never completes its designated tasks, a starvation situation arises.

Starvation is a phenomenon where a process

keeps getting more and more priority than other processes.

It happens due to the fact that the process

is taking too much time to complete its task.

As a result, other processes are not able to